



screenshot

Cost Comparison



Cost Savings



VALUE

SAVINGS

\$627.46

SAVINGS %

+85.3%

TIMELINE

BREAK-EVEN

87K reqs

EARLIEST SWITCH

2.8 min

TRADE-OFF

	ACCURACY	ITEMS / S	TTL.
LLM	0.93	13	213.7 h
Poodle	0.91	254	11.0 h

Model Prices

MODEL	INPUT \$/1M	OUTPUT \$/1M	PROVIDER	
BERT 80M small	\$0.01	\$0.01	Together AI	
Llama 3.1 8B large	\$0.20	\$0.20	Together AI	
Gemini 2.5 Flash Lite	\$0.10	\$0.40	Google	
Gemini 2.5 Flash	\$0.30	\$2.50	Google	
Gemini 2.5 Pro	\$1.25	\$10.00	Google	
Gemini 3 Flash Preview	\$0.50	\$3.00	Google	
Model name	Provider	Input \$/1M	Output \$/1M	+ Add
Claude Sonnet 5	\$3.00	\$15.00	Anthropic	

Use Case

imdb

Input

"A Mouse in the House" is a very classic cartoon by Tom & Jerry, faithful to their tradition but with jokes of its own. It is hysterical, hilarious, very entertaining and quite amusing. Artwork is of good quality either. This short isn't just about Tom trying to catch Jerry. Butch lives in the same house and he's trying to catch the mouse too, because <the's only going to be one cat in this house in the morning - and that's the cat that catches the mouse>.

If you ask me, there are lots of funny gags in this cartoon. The funniest for me are, for example, when Mammy Two Shoes sees the two lazy cats

Prompt

You are a sentiment analysis assistant. Your job is to read movie reviews and classify their sentiment as either 'positive' or 'negative'. Only respond in this exact JSON format: {"sentiment": "positive"} or {"sentiment": "negative"}. Do not provide any explanation or

Output

{"sentiment": "positive"}

Models

Large Model

Llama 3.1 8B

Small Model

BERT 80M

Requests

Expected Requests

10,000,000

100000

Switch After N Requests

5,000

5000

Model Development

s-500

Details →

Model Dev Costs (\$)

\$5.00

5

Validation

Validation Requests Percent

9%

0,09

Task Detection

User Provided

Wrapped Requests

100%

1

Wrapper Prompt

You are a powerful language model. You are given a user request consisting of a system prompt and a user message. Your task is as follows:

1. Identify the type of input the user is providing: one of ["text", "image", "table", "other"]
2. Infer what task you are expected to perform, choosing from:

Wrapped Output

{
 "input_type": "text",
 "task_type": "sentiment classification",
 "user_response": {
 "sentiment": "positive"
 }
}